

Assignment-6

GitHub link: <https://github.com/SUSHMASHETTY6/Assignment-6>

Question 1: Mathematical Solution


Q1) Calculate and find out clustering representation and dendrogram using single, complete, and average link proximity function in hierarchical clustering technique.

Point	Xcoordinate	Ycoordinate
P1	0.4005	0.5306
P2	0.2148	0.3854
P3	0.3457	0.3156
P4	0.2652	0.1875
P5	0.0789	0.4139
P6	0.4548	0.3022

	P1	P2	P3	P4	P5	P6
P1	0.0000	0.2357	0.2218	0.3688	0.3421	0.2347
P2	0.2357	0.0000	0.1483	0.2042	0.1388	0.2540
P3	0.2218	0.1483	0.0000	0.1513	0.2843	0.1180
P4	0.3688	0.2042	0.1513	0.0000	0.2932	0.2116
P5	0.3421	0.1388	0.2843	0.2932	0.0000	0.3921
P6	0.2347	0.2540	0.1180	0.2116	0.3921	0.0000

By single link :-

- * For single link hierarchical clustering, the proximity of two clusters is minimum of the distance between any two points in 2 different clusters.
- * The single link technique is good for non elliptical shapes, but sensitive to noise & outliers.
- * Applying single link technique to our example data set of six points.



→ from table 1, we can observe distance between P5 & P6 is 0.11


→ The height at which two clusters are merged can be represented as distance between two clusters.

Distance between clusters {5, 6} & {2, 4} is given by


$$\text{dist} \{3, 6\}, \{2, 5\} = \min(\text{dist}(3, 2), \text{dist}(6, 2), \text{dist}(3, 5), \text{dist}(6, 5))$$

$$\Rightarrow \min(0.15, 0.25, 0.28, 0.39)$$

$$\Rightarrow 0.15$$



single link clustering



single link dendrogram

Complete link :-

→ In Complete link of hierarchical clustering, the point of two clusters is defined as "the maximum of the distance between any two points in two different clusters."

→ Complete link is less susceptible to noise & outliers, but it can break large clusters & it favours globular shapes.


→ Below figure shows results of applying max to the sample data set of six points.

→ Points 3 & 6 are merged first.
 $\{3, 6\}$ is merged with $\{4\}$ instead of $\{2, 5\}$ or $\{1\}$ this is because


$$\begin{aligned} \text{dist}(\{3, 6\}, \{4\}) &= \max(\text{dist}(3, 4), \text{dist}(6, 4)) \\ &= \max(0.15, 0.22) \\ &= 0.22 \end{aligned}$$

$$\begin{aligned} \text{dist}(\{3, 6\}, \{2, 5\}) &= \max(\text{dist}(3, 2), \text{dist}(6, 2), \text{dist}(3, 5), \text{dist}(6, 5)) \\ &= \max(0.15, 0.25, 0.28, 0.39) \\ &= 0.39 \end{aligned}$$

$$\begin{aligned} \text{dist}(\{3, 6\}, \{1\}) &= \max(\text{dist}(3, 1), \text{dist}(6, 1)) \\ &= \max(0.22, 0.23) \\ &= 0.23 \end{aligned}$$



complete link dendrogram



complete link dendrogram

Average link:-

Below fig shows results after applying the group average approach to sample data of five points.

→ we calculate the distance between some cluster


$$\rightarrow \text{proximity} \Rightarrow \text{proximity}(i, j) = \frac{\sum_{x \in i} \sum_{y \in j} \text{proximity}(x, y)}{n_i n_j}$$

$$\text{dist}(\{3, 6, 4\}, \{1\}) = (0.22 + 0.37 + 0.23) / (3 \times 1) = 0.28$$

$$\text{dist}(\{2, 5\}, \{1\}) = (0.24 + 0.34) / (2 \times 1) = 0.29$$

$$\text{dist}(\{3, 6, 4\}, \{2, 5\}) = (0.15 + 0.28 + 0.25 + 0.39 + 0.20 + 0.23) / 6 = 0.26$$

Here, Because $\text{dist}(\{3, 6, 4\}, \{2, 5\})$ is smaller than $\text{dist}(\{3, 6, 4\}, \{1\})$ and $\text{dist}(\{2, 5\}, \{1\})$ clusters $\{3, 6, 4\}$ and $\{2, 5\}$ are merged at the fourth stage.



1) group average dendrogram

→ Average version of hierarchical clustering, the proximity of two clusters is defined as the average pairwise proximity among all pairs of points in the different clusters.

proximity proximity (G, G') of clusters G and G' which are of size m_i and m_j respectively

$$\text{Proximity } (G, G') = \frac{\sum_{x \in G, y \in G'} \text{Proximity}(x, y)}{m_i \times m_j}$$

Question 2: Screenshots:

```
In [52]: #importing all Libraries here for assignment
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, metrics
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.decomposition import PCA
from sklearn.cluster import AgglomerativeClustering
from sklearn.metrics import silhouette_score

import warnings
warnings.filterwarnings("ignore")

In [53]: dataframe = pd.read_csv('CC General.csv')
dataframe.info()
```

Reads the csv file

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8950 entries, 0 to 8949
Data columns (total 18 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   CUST_ID                               8950 non-null   object
 1   BALANCE                               8950 non-null   float64
 2   BALANCE_FREQUENCY                     8950 non-null   float64
 3   PURCHASES                             8950 non-null   float64
 4   ONEOFF_PURCHASES                      8950 non-null   float64
 5   INSTALLMENTS_PURCHASES                8950 non-null   float64
 6   CASH_ADVANCE                          8950 non-null   float64
 7   PURCHASES_FREQUENCY                   8950 non-null   float64
 8   ONEOFF_PURCHASES_FREQUENCY            8950 non-null   float64
 9   PURCHASES_INSTALLMENTS_FREQUENCY      8950 non-null   float64
10  CASH_ADVANCE_FREQUENCY                8950 non-null   float64
11  CASH_ADVANCE_TRX                      8950 non-null   int64
12  PURCHASES_TRX                         8950 non-null   int64
13  CREDIT_LIMIT                          8949 non-null   float64
14  PAYMENTS                              8950 non-null   float64
15  MINIMUM_PAYMENTS                      8637 non-null   float64
16  PRC_FULL_PAYMENT                      8950 non-null   float64
17  TENURE                                8950 non-null   int64
dtypes: float64(14), int64(3), object(1)
memory usage: 1.2+ MB
```

dataframe.head(): This function returns the first n rows for the object based on position.

```
In [54]: dataframe.head()

Out[54]:
```

	CUST_ID	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLMENTS_PURCHASES	CASH_ADVANCE	PURCHASES_FREQUENCY
0	C10001	40.900749	0.818182	95.40	0.00	95.4	0.000000	
1	C10002	3202.467416	0.909091	0.00	0.00	0.0	6442.945483	
2	C10003	2495.148862	1.000000	773.17	773.17	0.0	0.000000	
3	C10004	1666.670542	0.636364	1499.00	1499.00	0.0	205.788017	
4	C10005	817.714335	1.000000	16.00	16.00	0.0	0.000000	

The describe() method returns description of the data in the DataFrame.

If the DataFrame contains numerical data, the description contains this information for each column.

```
In [55]: dataframe.describe()
```

```
Out[55]:
```

	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLMENTS_PURCHASES	CASH_ADVANCE	PURCHASES_FREQUENCY
count	8950.000000	8950.000000	8950.000000	8950.000000	8950.000000	8950.000000	8950.000000
mean	1564.474828	0.877271	1003.204834	592.437371	411.067645	978.871112	0.490000
std	2081.531879	0.236904	2136.634782	1659.887917	904.338115	2097.163877	0.401000
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	128.281915	0.888889	39.635000	0.000000	0.000000	0.000000	0.083333
50%	873.385231	1.000000	361.280000	38.000000	89.000000	0.000000	0.500000
75%	2054.140036	1.000000	1110.130000	577.405000	468.637500	1113.821139	0.916667
max	19043.138560	1.000000	49039.570000	40761.250000	22500.000000	47137.211760	1.000000

Remove rows or columns by specifying label names and corresponding axis(in this case its 1)

```
In [56]: df = dataframe.drop(['CUST_ID'], axis=1)
df.head()
```

```
Out[56]:
```

	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLMENTS_PURCHASES	CASH_ADVANCE	PURCHASES_FREQUENCY
0	40.900749	0.818182	95.40	0.00	95.4	0.000000	0.166667
1	3202.467416	0.909091	0.00	0.00	0.0	6442.945483	0.000000
2	2495.148862	1.000000	773.17	773.17	0.0	0.000000	1.000000
3	1666.670542	0.636364	1499.00	1499.00	0.0	205.788017	0.083333
4	817.714335	1.000000	16.00	16.00	0.0	0.000000	0.083333

returns Boolean value


```
In [57]: df.isnull().any()
```

```
Out[57]: BALANCE                False
          BALANCE_FREQUENCY      False
          PURCHASES              False
          ONEOFF_PURCHASES       False
          INSTALLMENTS_PURCHASES False
          CASH_ADVANCE           False
          PURCHASES_FREQUENCY    False
          ONEOFF_PURCHASES_FREQUENCY False
          PURCHASES_INSTALLMENTS_FREQUENCY False
          CASH_ADVANCE_FREQUENCY False
          CASH_ADVANCE_TRX       False
          PURCHASES_TRX         False
          CREDIT_LIMIT          True
          PAYMENTS              False
          MINIMUM_PAYMENTS      True
          PRC_FULL_PAYMENT      False
          TENURE                False
          dtype: bool
```

modifies if any values are true

```
In [58]: df.fillna(dataframe.mean(), inplace=True)
          df.isnull().any()
```

```
Out[58]: BALANCE                False
          BALANCE_FREQUENCY      False
          PURCHASES              False
          ONEOFF_PURCHASES       False
          INSTALLMENTS_PURCHASES False
          CASH_ADVANCE           False
          PURCHASES_FREQUENCY    False
          ONEOFF_PURCHASES_FREQUENCY False
          PURCHASES_INSTALLMENTS_FREQUENCY False
          CASH_ADVANCE_FREQUENCY False
          CASH_ADVANCE_TRX       False
          PURCHASES_TRX         False
          CREDIT_LIMIT          False
          PAYMENTS              False
          MINIMUM_PAYMENTS      False
          PRC_FULL_PAYMENT      False
          TENURE                False
          dtype: bool
```

color grid variances

```
In [59]: df.corr().style.background_gradient(cmap="Greens")
```

	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLMENTS_PURCHASES	CASH_A
BALANCE	1.000000	0.322412	0.181261	0.164350	0.126469	
BALANCE_FREQUENCY	0.322412	1.000000	0.133674	0.104323	0.124292	
PURCHASES	0.181261	0.133674	1.000000	0.916845	0.679896	
ONEOFF_PURCHASES	0.164350	0.104323	0.916845	1.000000	0.330622	
INSTALLMENTS_PURCHASES	0.126469	0.124292	0.679896	0.330622	1.000000	
CASH_ADVANCE	0.496692	0.099388	-0.051474	-0.031326	-0.064244	
PURCHASES_FREQUENCY	-0.077944	0.229715	0.393017	0.264937	0.442418	
ONEOFF_PURCHASES_FREQUENCY	0.073166	0.202415	0.498430	0.524891	0.214042	
PURCHASES_INSTALLMENTS_FREQUENCY	-0.063186	0.176079	0.315567	0.127729	0.511351	
CASH_ADVANCE_FREQUENCY	0.449218	0.191873	-0.120143	-0.082628	-0.132318	
CASH_ADVANCE_TRX	0.385152	0.141555	-0.067175	-0.046212	-0.073999	
PURCHASES_TRX	0.154338	0.189626	0.689561	0.545523	0.628108	
CREDIT_LIMIT	0.531267	0.095795	0.356959	0.319721	0.256496	
PAYMENTS	0.322802	0.065008	0.603264	0.567292	0.384084	
MINIMUM_PAYMENTS	0.394282	0.114249	0.093515	0.048597	0.131687	
PRC_FULL_PAYMENT	-0.318959	-0.095082	0.180379	0.132763	0.182569	
TENURE	0.072692	0.119776	0.086288	0.064150	0.086143	

```
In [69]: x = df.iloc[:,0:-1]
y = df.iloc[:,1]

scaler = preprocessing.StandardScaler()
scaler.fit(x)
X_scaled_array = scaler.transform(x)
X_scaled_df = pd.DataFrame(X_scaled_array, columns = x.columns)
```

```
In [70]: #Normalization is the process of scaling individual samples to have unit norm.
#This process can be useful if you plan to use a quadratic form such as the dot-product or any other kernel to
X_normalized = preprocessing.normalize(X_scaled_df)
# Converting the numpy array into a pandas DataFrame
X_normalized = pd.DataFrame(X_normalized)
```

Dataframe → principal components(P1,P2,Tenure)

```
In [62]: pca2 = PCA(n_components=2)
principalComponents = pca2.fit_transform(X_normalized)

principalDf = pd.DataFrame(data = principalComponents, columns = ['P1', 'P2'])


finalDf = pd.concat([principalDf, df[['TENURE']]], axis = 1)
finalDf.head()
```

Out[62]:

	P1	P2	TENURE
0	-0.488186	-0.677233	12
1	-0.517294	0.556075	12
2	0.334384	0.287312	12
3	-0.486616	-0.080780	12
4	-0.562175	-0.474770	12

```
In [63]: ▶ plt.figure(figsize=(7,7))
plt.scatter(finalDf['P1'],finalDf['P2'],c=finalDf['TENURE'],cmap='prism', s =5)
plt.xlabel('pc1')
plt.ylabel('pc2')
```


Out[63]: Text(0, 0.5, 'pc2')



Agglomerative clustering(n=2)

```
In [64]: ▶ ac2 = AgglomerativeClustering(n_clusters = 2)


# Visualizing the clustering
plt.figure(figsize =(6, 6))
plt.scatter(principalDf['P1'], principalDf['P2'],
            c = ac2.fit_predict(principalDf), cmap ='rainbow')
plt.show()
```



Agglomerative clustering(n=3)

```
In [65]: ► ac3 = AgglomerativeClustering(n_clusters = 3)


# Visualizing the clustering
plt.figure(figsize =(6, 6))
plt.scatter(principalDf['P1'], principalDf['P2'],
            c = ac3.fit_predict(principalDf), cmap ='rainbow')
plt.show()
```



Agglomerative clustering(n=4)

```
In [66]: ▶ ac4 = AgglomerativeClustering(n_clusters = 4)


# Visualizing the clustering
plt.figure(figsize =(6, 6))
plt.scatter(principalDf['P1'], principalDf['P2'],
            c = ac4.fit_predict(principalDf), cmap ='rainbow')
plt.show()
```

Agglomerative clustering(n=5)

```
In [67]: ► ac5 = AgglomerativeClustering(n_clusters = 5)

# Visualizing the clustering
plt.figure(figsize =(6, 6))
plt.scatter(principalDf['P1'], principalDf['P2'],
            c = ac5.fit_predict(principalDf), cmap ='rainbow')
plt.show()
```



Appending the silhouette score of the different models to the list and then plotting a bar graph to compare the results

```
In [68]: k = [2, 3, 4, 5]

# Appending the silhouette scores of the different models to the list
silhouette_scores = []
silhouette_scores.append(
    silhouette_score(principalDf, ac2.fit_predict(principalDf)))
silhouette_scores.append(
    silhouette_score(principalDf, ac3.fit_predict(principalDf)))
silhouette_scores.append(
    silhouette_score(principalDf, ac4.fit_predict(principalDf)))
silhouette_scores.append(
    silhouette_score(principalDf, ac5.fit_predict(principalDf)))

# Plotting a bar graph to compare the results
plt.bar(k, silhouette_scores)
plt.xlabel('Number of clusters', fontsize = 20)
plt.ylabel('S(i)', fontsize = 20)
plt.show()
```

